

Tying the (Data) Knot: Love, Life & Likes

WIA1006/WID3006 Machine Learning - Group Assignment

Objective: Predict relationship match outcomes from a synthetic dating app dataset (50,000 records) using multiple ML models.

Group 8 Members:

- Chan Jun Foong (Leader) - Pipeline Architecture & XGBoost
- Lim Wei Jie - Data Preprocessing & Feature Engineering
- Nurul Aisyah - Logistic Regression & Feature Selection
- Muhammad Irfan - Random Forest & Hyperparameter Tuning
- Tan Shin Yee - Decision Tree & Error Analysis
- Priya Devi - SVM & Ensembling

Table of Contents

1. [Task 1: Data Preprocessing](#)
2. [Task 2: Feature Engineering & Selection](#)
3. [Task 3: Model Selection & Training](#)
4. [Task 4: Model Evaluation & Comparison](#)
5. [Task 5: Hyperparameter Tuning](#)
6. [Bonus: Balancing Techniques \(SMOTE\)](#)
7. [Conclusion](#)

Setup: Install & Import Dependencies

In [1]:

```
# Install required packages (uncomment if needed) # !pip install pandas numpy matplotlib seaborn scikit-learn xgboost shap
imbalanced-learn joblib import warnings warnings.filterwarnings("ignore") import os import sys import numpy as np import pandas as pd
import matplotlib.pyplot as plt import seaborn as sns import joblib from sklearn.model_selection import train_test_split,
RandomizedSearchCV, learning_curve, cross_val_score from sklearn.pipeline import Pipeline from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder, label_binarize from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectKBest, f_classif from sklearn.metrics import (accuracy_score, precision_score, recall_score,
f1_score, roc_auc_score, classification_report, confusion_matrix, ConfusionMatrixDisplay, roc_curve, auc) from sklearn.linear_model
import LogisticRegression from sklearn.tree import DecisionTreeClassifier from sklearn.ensemble import RandomForestClassifier,
VotingClassifier from sklearn.svm import LinearSVC from sklearn.neural_network import MLPClassifier from xgboost import XGBClassifier
from imblearn.pipeline import Pipeline as ImbPipeline from imblearn.combine import SMOTEENN # Visualization settings
sns.set_theme(style="whitegrid", font_scale=1.1) plt.rcParams['figure.dpi'] = 150 PALETTE = "viridis" # Paths BASE_DIR =
os.path.abspath(os.path.join(os.getcwd(), "..")) DATA_PATH = os.path.join(BASE_DIR, "data", "dating_app_behavior_dataset.csv") MODELS_DIR
= os.path.join(BASE_DIR, "models") VIZ_DIR = os.path.join(BASE_DIR, "visualizations") REPORTS_DIR = os.path.join(BASE_DIR, "reports") for
d in [MODELS_DIR, VIZ_DIR, REPORTS_DIR]: os.makedirs(d, exist_ok=True) print("Setup complete.")
```

Setup complete.

1. Task 1: Data Preprocessing

1.1 Load the Dataset

In [2]:

```
df = pd.read_csv(DATA_PATH) print(f"Dataset shape: {df.shape[0]} rows x {df.shape[1]} columns") print(f"\nColumn types:")
print(df.dtypes) df.head()
```

```
Dataset shape: 50000 rows x 19 columns Column types: gender str sexual_orientation str location_type str income_bracket str
education_level str interest_tags str app_usage_time_min int64 app_usage_time_label str swipe_right_ratio float64 swipe_right_label str
likes_received int64 mutual_matches int64 profile_pics_count int64 bio_length int64 message_sent_count int64 emoji_usage_rate float64
last_active_hour int64 swipe_time_of_day str match_outcome str dtype: object
```

Out[2]:

	gender	sexual_orientat	ocation_type	come_brace	education_level	interest_apps	usage_tapp	usage_time	swipe_right	swipe_right_likes	received_mutual_matches	profile_pics	bio_length	message_sent	emoji_usage	last_active	swipe_time_of_day	match_outcome	
0	Prefer Not to Say	Gay	Urban	High	Bachelor's	Fitness, Politics, Traveling	52	Moderate	0.60	Optimistic	173	23	4	44	75	0.36	13	Early Morning	Mutual Match
1	Male	Bisexual	Suburban	Upper-Middle	No Formal Education	Languages, Fashion, Parenting	279	Extreme User	0.56	Optimistic	107	7	3	301	35	0.42	0	Morning	Chat Ignored
2	Non-binary	Pansexual	Suburban	Low	Master's	Movies, Reading, DIY	49	Moderate	0.41	Optimistic	91	27	2	309	33	0.41	1	After Midnight	Date Happened
3	Genderfluid	Gay	Metro	Very Low	Postdoc	Coding, Podcasts, History	185	Extreme User	0.32	Balanced	147	6	5	35	5	0.07	21	Morning	No Action
4	Male	Bisexual	Urban	Middle	Bachelor's	Clubbing, Podcasts, Cars	83	High	0.32	Balanced	94	11	1	343	34	0.11	22	After Midnight	One-sided Like

1.2 Summary Statistics

In [3]:

```
# Summary statistics for numerical features df.describe()
```

Out[3]:

	app_usage_time_min	swipe_right_ratio	likes_received	mutual_matches	profile_pics_count	bio_length	message_sent_count	emoji_usage_rate	last_active_hour
count	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000
mean	149.912400	0.500655	99.526040	13.870280	2.98772	250.174400	50.07194	0.286205	11.521800
std	86.990521	0.197468	57.996799	9.105615	1.99678	144.800996	29.16800	0.160042	6.920474
min	0.000000	0.000000	0.000000	0.000000	0.00000	0.000000	0.00000	0.000000	0.000000
25%	74.000000	0.370000	49.000000	6.000000	1.00000	125.000000	25.00000	0.160000	5.000000
50%	150.000000	0.500000	100.000000	13.000000	3.00000	250.000000	50.00000	0.270000	12.000000
75%	225.000000	0.640000	150.000000	22.000000	5.00000	376.000000	75.00000	0.390000	18.000000
max	300.000000	1.000000	200.000000	30.000000	6.00000	500.000000	100.00000	0.940000	23.000000

In [4]:

```
# Summary statistics for categorical features df.describe(include=['object'])
```

Out[4]:

	gender	sexual_orientation	location_type	income_bracket	education_level	interest_tags	app_usage_time_label	swipe_right_label	swipe_time_of_day	match_outcome
count	50000	50000	50000	50000	50000	50000	50000	50000	50000	50000
unique	6	8	6	7	9	40206	7	4	6	10
top	Female	Straight	Remote Area	High	Bachelor's	Fitness, Anime, Yoga	Extreme User	Optimistic	After Midnight	One-sided Like
freq	8384	6326	8519	7309	5646	6	20140	26873	8524	5112

1.3 Data Cleaning

Dropped columns (justification):

- mutual_matches: Target leakage - directly encodes match information
- app_usage_time_label: Derived bin of app_usage_time_min (redundant)
- swipe_right_label: Derived bin of swipe_right_ratio (redundant)
- sexual_orientation: Not in recommended feature list

In [5]:

```
DROP_COLS = [ "mutual_matches", # Target leakage "app_usage_time_label", # Redundant (derived bin) "swipe_right_label", # Redundant (derived bin) "sexual_orientation", # Not recommended ] existing_drops = [c for c in DROP_COLS if c in df.columns] df = df.drop(columns=existing_drops) print(f"Dropped columns: {existing_drops}") print(f"Shape after dropping: {df.shape}")

Dropped columns: ['mutual_matches', 'app_usage_time_label', 'swipe_right_label', 'sexual_orientation'] Shape after dropping: (50000, 15)
```

1.4 Target Variable: Group 10 Outcomes into 3 Classes

The original match_outcome has 10 categories. We group them into 3 meaningful classes:

- **Positive:** Mutual Match, Date Happened, Relationship Formed, Instant Match
- **Neutral:** No Action, One-sided Like, Chat Ignored
- **Negative:** Ghosted, Blocked, Catfished

In [6]:

```
TARGET = "match_outcome" TARGET_MAP = { "Mutual Match": "Positive", "Date Happened": "Positive", "Relationship Formed": "Positive",
"Instant Match": "Positive", "No Action": "Neutral", "One-sided Like": "Neutral", "Chat Ignored": "Neutral", "Ghosted": "Negative",
"Blocked": "Negative", "Catfished": "Negative", } df[TARGET] = df[TARGET].map(TARGET_MAP) unmapped = df[TARGET].isna().sum() if unmapped:
print(f"WARNING: {unmapped} unmapped rows - dropping") df = df.dropna(subset=[TARGET]) print("Target distribution after grouping:")
print(df[TARGET].value_counts()) # Encode target label_encoder = LabelEncoder() df[TARGET] = label_encoder.fit_transform(df[TARGET])
CLASS_NAMES = list(label_encoder.classes_) print(f"\nEncoded classes: {CLASS_NAMES}")
```

```
Target distribution after grouping: match_outcome Positive 19850 Neutral 15117 Negative 15033 Name: count, dtype: int64 Encoded classes:
['Negative', 'Neutral', 'Positive']
```

1.5 Handle Missing Values & Outliers

In [7]:

```
# Check for missing values missing = df.isnull().sum() print("Missing values per column:") print(missing[missing > 0] if missing.sum() >
0 else "No missing values found.")
```

```
Missing values per column: No missing values found.
```

In [8]:

```
# Outlier detection using IQR method numerical_cols_check = ['app_usage_time_min', 'swipe_right_ratio', 'likes_received',
'profile_pics_count', 'bio_length', 'message_sent_count', 'emoji_usage_rate', 'last_active_hour'] outlier_summary = {} for col in
numerical_cols_check: Q1 = df[col].quantile(0.25) Q3 = df[col].quantile(0.75) IQR = Q3 - Q1 lower = Q1 - 1.5 * IQR upper = Q3 + 1.5 * IQR
outliers = ((df[col] < lower) | (df[col] > upper)).sum() outlier_summary[col] = outliers outlier_df =
pd.DataFrame.from_dict(outlier_summary, orient='index', columns=['Outlier Count']) print("Outlier counts (IQR method):")
print(outlier_df)
```

```
Outlier counts (IQR method): Outlier Count app_usage_time_min 0 swipe_right_ratio 0 likes_received 0 profile_pics_count 0 bio_length 0
message_sent_count 0 emoji_usage_rate 311 last_active_hour 0
```

1.6 Categorical Encoding & Numerical Scaling

We build a `ColumnTransformer` pipeline:

- **Numerical:** `SimpleImputer (median) → StandardScaler`
- **Categorical:** `SimpleImputer (mode) → OneHotEncoder`

In [9]:

```
NUMERICAL_FEATURES = [ "app_usage_time_min", "swipe_right_ratio", "likes_received", "profile_pics_count", "bio_length",
"message_sent_count", "emoji_usage_rate", "last_active_hour", ] CATEGORICAL_FEATURES = [ "gender", "education_level", "income_bracket",
"location_type", "swipe_time_of_day", ] print(f"Numerical features ({len(NUMERICAL_FEATURES)}): {NUMERICAL_FEATURES}")
print(f"Categorical features ({len(CATEGORICAL_FEATURES)}): {CATEGORICAL_FEATURES}")
```

```
Numerical features (8): ['app_usage_time_min', 'swipe_right_ratio', 'likes_received', 'profile_pics_count', 'bio_length',
'message_sent_count', 'emoji_usage_rate', 'last_active_hour'] Categorical features (5): ['gender', 'education_level', 'income_bracket',
'location_type', 'swipe_time_of_day']
```

1.7 Exploratory Data Analysis (EDA)


In [10]:

```
# 1. Class Distribution fig, ax = plt.subplots(figsize=(8, 5)) counts = df[TARGET].value_counts().sort_index() labels = [CLASS_NAMES[i]
for i in counts.index] bars = ax.bar(labels, counts.values, color=sns.color_palette(PALETTE, len(counts)), edgecolor="black")
ax.set_title("Target Class Distribution", fontsize=14, fontweight="bold") ax.set_xlabel("Match Outcome Group") ax.set_ylabel("Count") for
bar, v in zip(bars, counts.values): ax.text(bar.get_x() + bar.get_width()/2, v + 100, f"{v:,}", ha="center", fontsize=10)
plt.tight_layout() plt.savefig(os.path.join(VIZ_DIR, "01_class_distribution.png"), dpi=150, bbox_inches="tight") plt.show()
```



In [11]:


```
# 2. Correlation Heatmap fig, ax = plt.subplots(figsize=(12, 10)) corr = df[NUMERICAL_FEATURES].corr() mask = np.triu(np.ones_like(corr, dtype=bool)) sns.heatmap(corr, mask=mask, annot=True, fmt=".2f", cmap="coolwarm", center=0, linewidths=0.5, ax=ax) ax.set_title("Pearson Correlation Heatmap", fontsize=14, fontweight="bold") plt.tight_layout() plt.savefig(os.path.join(VIZ_DIR, "02_correlation_heatmap.png"), dpi=150, bbox_inches="tight") plt.show()
```



In [12]:

```
# 3. Boxplots for Outlier Detection
n = len(NUMERICAL_FEATURES)
cols = 3
rows = (n + cols - 1) // cols
fig, axes = plt.subplots(rows, cols, figsize=(5 * cols, 4 * rows))
axes = axes.flatten()
for i, col in enumerate(NUMERICAL_FEATURES):
    sns.boxplot(y=df[col], ax=axes[i], color=sns.color_palette(PALETTE, n)[i])
    axes[i].set_title(col, fontsize=10)
    for j in range(i + 1, len(axes)):
        axes[j].set_visible(False)
fig.suptitle("Outlier Detection - Boxplots", fontsize=14, fontweight="bold", y=1.02)
fig.tight_layout()
plt.savefig(os.path.join(VIZ_DIR, "03_boxplots.png"), dpi=150, bbox_inches="tight")
plt.show()
```


Outlier Detection - Boxplots



In [13]:

```
# 4. Feature Distributions by Target Class selected = NUMERICAL_FEATURES[:6] fig, axes = plt.subplots(2, 3, figsize=(16, 10)) axes =
axes.flatten() target_labels = {i: name for i, name in enumerate(CLASS_NAMES)} for i, col in enumerate(selected): for cls_val in
sorted(df[TARGET].unique()): subset = df[df[TARGET] == cls_val][col] axes[i].hist(subset, bins=30, alpha=0.5,
label=target_labels[cls_val]) axes[i].set_title(col) axes[i].legend(fontsize=7) fig.suptitle("Feature Distributions by Match Outcome",
fontsize=14, fontweight="bold") fig.tight_layout() plt.savefig(os.path.join(VIZ_DIR, "04_distributions_by_target.png"), dpi=150,
bbox_inches="tight") plt.show()
```

Feature Distributions by Match Outcome



2. Task 2: Feature Engineering & Selection

2.1 Create New Features

We engineer 5 composite features:

- **EngagementScore** = likes_received + message_sent_count (passive + active interaction)
- **ProfileQuality** = bio_length + 5 × profile_pics_count (profile effort metric)
- **ActivityIntensity** = app_usage_time_min × swipe_right_ratio (engagement × aggressiveness)
- **interest_tags_count** = count of comma-separated interest tags
- **EngagementIntensity** = emoji_usage_rate × message_sent_count (enthusiasm signal)

In [14]:

```
df["EngagementScore"] = df["likes_received"] + df["message_sent_count"] df["ProfileQuality"] = df["bio_length"] + 5 *
df["profile_pics_count"] df["ActivityIntensity"] = df["app_usage_time_min"] * df["swipe_right_ratio"] if "interest_tags" in df.columns:
df["interest_tags_count"] = df["interest_tags"].apply( lambda x: len(str(x).split(',')) if pd.notna(x) else 0 ) df["EngagementIntensity"]
= df["emoji_usage_rate"] * df["message_sent_count"] # Update feature lists ENGINEERED_FEATURES = ["EngagementScore", "ProfileQuality",
"ActivityIntensity", "interest_tags_count", "EngagementIntensity"] ALL_NUMERICAL = NUMERICAL_FEATURES + ENGINEERED_FEATURES
print(f"Engineered features added: {ENGINEERED_FEATURES}") print(f"Total numerical features: {len(ALL_NUMERICAL)}")
df[ENGINEERED_FEATURES].describe()
```

```
Engineered features added: ['EngagementScore', 'ProfileQuality', 'ActivityIntensity', 'interest_tags_count', 'EngagementIntensity'] Total
numerical features: 13
```

Out[14]:

	EngagementScore	ProfileQuality	ActivityIntensity	interest_tags_count	EngagementIntensity
count	50000.000000	50000.000000	50000.000000	50000.0	50000.000000
mean	149.597980	265.113000	75.062878	3.0	14.345578
std	64.890425	145.220599	55.454516	0.0	12.499390
min	0.000000	0.000000	0.000000	3.0	0.000000

2.5%	100.000000	140.000000	29.600000	3.0	4.590000
5.0%	149.000000	265.000000	65.100000	3.0	11.000000
7.5%	200.000000	391.000000	110.337500	3.0	20.790000
10.0%	300.000000	530.000000	300.000000	3.0	84.600000

2.2 Feature Selection using SelectKBest

We use `SelectKBest` with `f_classif` (ANOVA F-test) to select the top 15 features from the combined numerical + one-hot encoded categorical feature space. This reduces dimensionality and removes noisy features.

In [15]:

```
# Quick preprocessor to demonstrate feature selection
def build_preprocessor():
    numerical_pipeline = Pipeline([ ("imputer", SimpleImputer(strategy="median")), ("scaler", StandardScaler()), ])
    categorical_pipeline = Pipeline([ ("imputer", SimpleImputer(strategy="most_frequent")), ("encoder", OneHotEncoder(handle_unknown="ignore", sparse_output=False)), ])
    return ColumnTransformer([ ("num", numerical_pipeline, ALL_NUMERICAL), ("cat", categorical_pipeline, CATEGORICAL_FEATURES), ])

# Fit preprocessor to see all feature names
prep = build_preprocessor()
X_temp = df[ALL_NUMERICAL + CATEGORICAL_FEATURES]
X_temp_transformed = prep.fit_transform(X_temp)

# Get feature names after one-hot encoding
ohe = prep.named_transformers_["cat"].named_steps["encoder"]
cat_names = list(ohe.get_feature_names_out())
all_feature_names = ALL_NUMERICAL + cat_names

# Apply SelectKBest selector
selector = SelectKBest(score_func=f_classif, k=15)
y_temp = df[TARGET]
selector.fit(X_temp_transformed, y_temp)

# Show selected features and their scores
scores = pd.DataFrame({ "Feature": all_feature_names, "F_Score": selector.scores_, "p_value": selector.pvalues_, "Selected": selector.get_support() })
scores.sort_values("F_Score", ascending=False)

print("Top 15 Selected Features (by ANOVA F-score):")
print(scores[scores["Selected"]].to_string(index=False))
print(f"\nTotal features: {len(all_feature_names)} → Selected: {selector.get_support().sum()}")
```

Top 15 Selected Features (by ANOVA F-score):

Feature	F_Score	p_value	Selected
x4_Morning	3.287031	0.037373	True
x1_Associate's	2.959824	0.051837	True
x1_MBA	2.486905	0.083177	True
x2_Middle	2.267488	0.103583	True
x4_Early Morning	1.842800	0.158384	True
x2_Low	1.724476	0.178277	True
ActivityIntensity	1.722394	0.178649	True
x1_PhD	1.679832	0.186416	True
message_sent_count	1.540172	0.214354	True
profile_pics_count	1.508097	0.221341	True
x3_Remote Area	1.452619	0.233967	True
x2_Upper-Middle	1.430342	0.239237	True
x1_No Formal Education	1.338943	0.262132	True
swipe_right_ratio	1.298550	0.272936	True
x3_Suburban	1.265375	0.282142	True

Total features: 47 → Selected: 15

2.3 Feature Importance (from tree-based models)

After training, we will visualize feature importance from Random Forest and XGBoost (see Section 4).

3. Task 3: Model Selection & Training

3.1 Train-Test Split (80/20, Stratified)

In [16]:

```
FEATURE_COLS = ALL_NUMERICAL + CATEGORICAL_FEATURES
X = df[FEATURE_COLS]
y = df[TARGET]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

print(f"Training set: {X_train.shape[0]} samples")
print(f"Test set: {X_test.shape[0]} samples")
print(f"\nClass distribution (train):")
print(y_train.value_counts().sort_index().rename(lambda i: CLASS_NAMES[i]))
print(f"\nClass distribution (test):")
print(y_test.value_counts().sort_index().rename(lambda i: CLASS_NAMES[i]))
```

Training set: 40000 samples Test set: 10000 samples Class distribution (train):

match_outcome	count
Negative	12026
Neutral	12094
Positive	15880

Class distribution (test):

match_outcome	count
Negative	3007
Neutral	3023
Positive	3970

3.2 Define Models

We train **7 models** across different algorithmic families:

#	Model	Type	Rationale
1	Logistic Regression	Linear	Strong linear baseline
2	Decision Tree	Rule-based	Interpretable, non-linear
3	Random Forest	Bagging Ensemble	Variance reduction
4	XGBoost	Boosting Ensemble	State-of-the-art performance
5	SVM (LinearSVC)	Margin-based	High-dimensional boundaries
6	MLP Neural Network	Deep Learning	Non-linear interactions
7	Voting Classifier	Ensemble	Soft voting (LR + RF + XGB)

In [17]:

```
preprocessor = build_preprocessor() models_and_params = [ ("Logistic Regression", LogisticRegression(max_iter=1000,
class_weight="balanced", random_state=42), {"C": [0.01, 0.1, 1.0, 10.0, 100.0]}), ("Decision Tree",
DecisionTreeClassifier(class_weight="balanced", random_state=42), {"max_depth": [10, 20, 30, None], "min_samples_split": [2, 5, 10, 20],
"min_samples_leaf": [1, 5, 10]}), ("Random Forest", RandomForestClassifier(class_weight="balanced", random_state=42, n_jobs=1),
{"n_estimators": [100, 200, 300], "max_depth": [10, 20, 30, None], "min_samples_split": [2, 5, 10]}), ("XGBoost",
XGBClassifier(eval_metric="mlogloss", random_state=42, n_jobs=1), {"learning_rate": [0.01, 0.05, 0.1, 0.2], "max_depth": [3, 5, 7, 10],
"n_estimators": [100, 200, 300], "subsample": [0.8, 1.0]}), ("SVM", LinearSVC(class_weight="balanced", random_state=42, max_iter=1000),
{"C": [0.01, 0.1, 1.0, 10.0]}), ("Neural Network (MLP)", MLPClassifier(activation="relu", max_iter=500, random_state=42,
early_stopping=True, validation_fraction=0.1), {"hidden_layer_sizes": [(64, 32), (128, 64), (128, 64, 32)], "alpha": [0.0001, 0.001,
0.01], "learning_rate_init": [0.001, 0.01]}), ] print(f"Defined {len(models_and_params)} models for training.")
```

Defined 6 models for training.

3.3 Train All Models with Hyperparameter Tuning (Task 5: RandomizedSearchCV)

Each model is wrapped in a pipeline: Preprocessor → SelectKBest(k=15) → Model

Hyperparameter tuning uses RandomizedSearchCV with:

- `n_iter=10` random combinations
- `cv=3` fold cross-validation
- `scoring='f1_weighted'` (optimizes for balanced F1 across all 3 classes)

In [18]:

```
fitted_models = {} for name, estimator, param_grid in models_and_params: print(f"\n{' '*60}") print(f" Training: {name}")
print(f"{' '*60}") pipe = Pipeline([ ("preprocessor", preprocessor), ("selector", SelectKBest(score_func=f_classif, k=15)), ("model",
estimator, )] if param_grid: grid = RandomizedSearchCV( pipe, param_distributions={"model__{k}": v for k, v in param_grid.items()}),
n_iter=10, cv=3, scoring="f1_weighted", n_jobs=-1, verbose=0, random_state=42, ) grid.fit(X_train, y_train) best_pipe =
grid.best_estimator_ print(f" Best params: {grid.best_params_}") print(f" Best CV F1: {grid.best_score_.4f}") else: pipe.fit(X_train,
y_train) best_pipe = pipe fitted_models[name] = best_pipe # Voting Classifier print(f"\n{' '*60}") print(" Training: Voting Classifier")
print(f"{' '*60}") estimators = [ ("lr", fitted_models["Logistic Regression"].named_steps["model"]), ("rf", fitted_models["Random
Forest"].named_steps["model"]), ("xgb", fitted_models["XGBoost"].named_steps["model"]), ] voting_clf =
VotingClassifier(estimators=estimators, voting="soft") voting_pipe = Pipeline([ ("preprocessor", preprocessor), ("selector",
SelectKBest(score_func=f_classif, k=15)), ("model", voting_clf), ]) voting_pipe.fit(X_train, y_train) fitted_models["Voting Classifier"]
= voting_pipe print(f"\nAll {len(fitted_models)} models trained.")
```

===== Training: Logistic Regression
=====

Best params: {'model__C': 0.1} Best CV F1: 0.3317 ===== Training: Decision Tree
=====

Best params: {'model__min_samples_split': 5, 'model__min_samples_leaf': 1, 'model__max_depth': None} Best CV F1: 0.3370
===== Training: Random Forest
=====

Best params: {'model__n_estimators': 200, 'model__min_samples_split': 5, 'model__max_depth': 20} Best CV F1: 0.3416
===== Training: XGBoost
=====

Best params: {'model__subsample': 0.8, 'model__n_estimators': 300, 'model__max_depth': 10, 'model__learning_rate': 0.2} Best CV F1:
0.3381 ===== Training: SVM
=====

Best params: {'model__C': 10.0} Best CV F1: 0.2817 ===== Training: Neural Network
(MLP) =====

Best params: {'model__learning_rate_init': 0.01, 'model__hidden_layer_sizes': (128, 64, 32), 'model__alpha': 0.0001} Best CV F1: 0.2370
===== Training: Voting Classifier
=====

All 7 models trained.

4. Task 4: Model Evaluation & Comparison

4.1 Evaluate All Models

Metrics computed:

- **Accuracy:** Overall correctness
- **Precision (weighted):** Positive predictive value

- **Recall (weighted):** Sensitivity
- **F1 (weighted):** Harmonic mean of precision & recall
- **ROC-AUC (OvR):** Area under ROC curve (One-vs-Rest)

In [19]:

```
classes = sorted(y_test.unique().tolist()) results = [] for name, model in fitted_models.items(): try: y_pred = model.predict(X_test) acc = accuracy_score(y_test, y_pred) prec = precision_score(y_test, y_pred, average="weighted", zero_division=0) rec = recall_score(y_test, y_pred, average="weighted", zero_division=0) f1 = f1_score(y_test, y_pred, average="weighted", zero_division=0) try: y_prob = model.predict_proba(X_test) y_test_bin = label_binarize(y_test, classes=classes) roc = roc_auc_score(y_test_bin, y_prob, average="weighted", multi_class="ovr") except Exception: roc = np.nan results.append({ "Model": name, "Accuracy": round(acc, 4), "Precision": round(prec, 4), "Recall": round(rec, 4), "F1_Weighted": round(f1, 4), "ROC_AUC": round(roc, 4) if not np.isnan(roc) else "N/A", }) print(f"\n{' '*60}\n {name}\n{' '*60}") print(classification_report(y_test, y_pred, target_names=CLASS_NAMES, zero_division=0)) except Exception as e: print(f"\n Skipping {name}: {e}") # Build comparison table comparison_df = pd.DataFrame(results) comparison_df = comparison_df.sort_values("F1_Weighted", ascending=False).reset_index(drop=True) comparison_df.index += 1 comparison_df.index.name = "Rank" comparison_df.to_csv(os.path.join(REPORTS_DIR, "model_comparison.csv")) print("\nModel Comparison Table (sorted by F1 Weighted):") comparison_df
```

```
===== Logistic Regression
===== precision recall f1-score support Negative 0.29 0.35 0.32 3007 Neutral 0.30 0.30 0.30 3023 Positive 0.40 0.33 0.36 3970 accuracy 0.33 10000 macro avg 0.33 0.33 0.33 10000 weighted avg 0.34 0.33 0.33 10000
===== Decision Tree =====
precision recall f1-score support Negative 0.31 0.34 0.32 3007 Neutral 0.30 0.31 0.31 3023 Positive 0.39 0.36 0.37 3970 accuracy 0.34 10000 macro avg 0.33 0.33 0.33 10000 weighted avg 0.34 0.34 0.34 10000
```

```
===== Random Forest =====
precision recall f1-score support Negative 0.31 0.26 0.28 3007 Neutral 0.30 0.26 0.28 3023 Positive 0.40 0.50 0.45 3970 accuracy 0.35 10000 macro avg 0.34 0.34 0.34 10000 weighted avg 0.34 0.35 0.35 10000
```

```
===== XGBoost =====
precision recall f1-score support Negative 0.32 0.26 0.29 3007 Neutral 0.31 0.27 0.29 3023 Positive 0.40 0.50 0.45 3970 accuracy 0.36 10000 macro avg 0.34 0.34 0.34 10000 weighted avg 0.35 0.36 0.35 10000
===== SVM
===== precision recall f1-score support Negative 0.30 0.06 0.11 3007 Neutral 0.31 0.02 0.03 3023 Positive 0.40 0.92 0.56 3970 accuracy 0.39 10000 macro avg 0.34 0.33 0.23 10000 weighted avg 0.34 0.39 0.26 10000
===== Neural Network (MLP)
===== precision recall f1-score support Negative 0.00 0.00 0.00 3007 Neutral 0.00 0.00 0.00 3023 Positive 0.40 1.00 0.57 3970 accuracy 0.40 10000 macro avg 0.13 0.33 0.19 10000 weighted avg 0.16 0.40 0.23 10000
```

```
===== Voting Classifier
===== precision recall f1-score support Negative 0.31 0.25 0.28 3007 Neutral 0.31 0.26 0.28 3023 Positive 0.40 0.51 0.45 3970 accuracy 0.36 10000 macro avg 0.34 0.34 0.34 10000 weighted avg 0.35 0.36 0.35 10000 Model Comparison Table (sorted by F1 Weighted):
```

Out[19]:


	Model	Accuracy	Precision	Recall	F1_Weighted	ROC_AUC
1	XGBoost	0.3579	0.3494	0.3579	0.3503	0.5093
2	Voting Classifier	0.3553	0.3458	0.3553	0.3466	0.5096
3	Random Forest	0.3541	0.3449	0.3541	0.3460	0.5057
4	Decision Tree	0.3371	0.3404	0.3371	0.3382	0.4996
5	Logistic Regression	0.3293	0.3361	0.3293	0.3307	0.4994
6	SVM	0.3895	0.3417	0.3895	0.2609	N/A
7	Neural Network (MLP)	0.3968	0.1576	0.3968	0.2256	0.5013

4.2 Confusion Matrices

In [20]:

```
n = len(fitted_models) fig, axes = plt.subplots(2, (n + 1) // 2, figsize=(6 * ((n + 1) // 2), 10)) axes = axes.flatten() for i, (name, model) in enumerate(fitted_models.items()): y_pred = model.predict(X_test) cm = confusion_matrix(y_test, y_pred, labels=classes) disp = ConfusionMatrixDisplay(cm, display_labels=CLASS_NAMES) disp.plot(ax=axes[i], cmap="Blues", colorbar=False, values_format="d") axes[i].set_title(name, fontsize=10, fontweight="bold") for j in range(i + 1, len(axes)): axes[j].set_visible(False) fig.suptitle("Confusion Matrices", fontsize=14, fontweight="bold") fig.tight_layout() plt.savefig(os.path.join(VIZ_DIR, "05_confusion_matrices.png"), dpi=150, bbox_inches="tight") plt.show()
```


Confusion Matrices



4.3 ROC Curves (One-vs-Rest)

In [21]:


```
y_test_bin = label_binarize(y_test, classes=classes) n_classes = len(classes) fig, axes = plt.subplots(1, n_classes, figsize=(7 * n_classes, 6)) if n_classes == 1: axes = [axes] for cls_idx, cls_name in enumerate(CLASS_NAMES): ax = axes[cls_idx] for name, model in fitted_models.items(): try: y_prob = model.predict_proba(X_test) fpr, tpr, _ = roc_curve(y_test_bin[:, cls_idx], y_prob[:, cls_idx]) roc_auc = auc(fpr, tpr) ax.plot(fpr, tpr, label=f"{name} (AUC={roc_auc:.3f})" except Exception: continue ax.plot([0, 1], [0, 1], "k--", alpha=0.4) ax.set_title(f"ROC - {cls_name}", fontweight="bold") ax.set_xlabel("False Positive Rate") ax.set_ylabel("True Positive Rate") ax.legend(fontsize=7, loc="lower right") fig.suptitle("ROC Curves (One-vs-Rest)", fontsize=14, fontweight="bold") fig.tight_layout() plt.savefig(os.path.join(VIZ_DIR, "06_roc_curves.png"), dpi=150, bbox_inches="tight") plt.show()
```



4.4 Model Comparison Bar Chart

In [22]:

```
metrics = ["Accuracy", "Precision", "Recall", "F1_Weighted"] df_plot = comparison_df.copy() fig, ax = plt.subplots(figsize=(12, 6)) x = np.arange(len(df_plot)) width = 0.18 for i, m in enumerate(metrics): vals = pd.to_numeric(df_plot[m], errors="coerce") ax.bar(x + i * width, vals, width, label=m) ax.set_xticks(x + width * 1.5) ax.set_xticklabels(df_plot["Model"], rotation=20, ha="right") ax.set_ylabel("Score") ax.set_title("Model Performance Comparison", fontsize=14, fontweight="bold") ax.legend() ax.set_ylim(0, 1.05) fig.tight_layout() plt.savefig(os.path.join(VIZ_DIR, "07_model_comparison.png"), dpi=150, bbox_inches="tight") plt.show()
```




4.5 Feature Importance (Random Forest & XGBoost)


In [23]:

```
for model_name in ["Random Forest", "XGBoost"]: model = fitted_models.get(model_name) if model is None: continue clf =
model.named_steps["model"] importances = clf.feature_importances_ # Get feature names prep = model.named_steps["preprocessor"] try: ohe =
prep.named_transformers_["cat"].named_steps["encoder"] cat_names_local = list(ohe.get_feature_names_out()) except Exception:
cat_names_local = [] num_names = list(prepare.transformers_[0][2]) all_names = num_names + cat_names_local selector =
model.named_steps.get("selector") if selector is not None: mask = selector.get_support() all_names = [n for n, m in zip(all_names, mask)
if m] if len(all_names) != len(importances): all_names = [f"f{i}" for i in range(len(importances))] idx = np.argsort(importances)[-15:]
fig, ax = plt.subplots(figsize=(8, 6)) ax.barh(range(len(idx)), importances[idx], color=sns.color_palette(PALETTE, len(idx)))
ax.set_yticks(range(len(idx))) ax.set_yticklabels([all_names[i] for i in idx]) ax.set_title(f"Feature Importance - {model_name}",
fontsize=14, fontweight="bold") ax.set_xlabel("Importance") fig.tight_layout() fname =
f"08_feature_importance_{model_name.lower().replace(' ', '_')}.png" plt.savefig(os.path.join(VIZ_DIR, fname), dpi=150,
bbox_inches="tight") plt.show()
```

Feature Importance - Random Forest




Feature Importance - XGBoost



4.6 Learning Curves (Overfitting Analysis)

In [24]:

```
selected_models = ["Logistic Regression", "Random Forest", "XGBoost"]
fig, axes = plt.subplots(1, len(selected_models), figsize=(7 * len(selected_models), 5))
for i, name in enumerate(selected_models):
    model = fitted_models.get(name)
    if model is None:
        continue
    train_sizes, train_scores, val_scores = learning_curve(
        model, X_train, y_train, cv=5, scoring="f1_weighted",
        train_sizes=np.linspace(0.1, 1.0, 8), n_jobs=-1,
    )
    axes[i].plot(train_sizes, train_scores.mean(axis=1), "o-", label="Train")
    axes[i].plot(train_sizes, val_scores.mean(axis=1), "o-", label="Validation")
    axes[i].fill_between(
        train_sizes,
        train_scores.mean(axis=1) - train_scores.std(axis=1),
        train_scores.mean(axis=1) + train_scores.std(axis=1),
        alpha=0.1,
    )
    axes[i].fill_between(
        train_sizes,
        val_scores.mean(axis=1) - val_scores.std(axis=1),
        val_scores.mean(axis=1) + val_scores.std(axis=1),
        alpha=0.1,
    )
    axes[i].set_title(f"Learning Curve - {name}")
    axes[i].set_xlabel("Training Samples")
    axes[i].set_ylabel("F1 (weighted)")
    axes[i].legend()
fig.suptitle("Learning Curves - Overfitting Analysis", fontsize=14, fontweight="bold")
fig.tight_layout()
plt.savefig(os.path.join(VIZ_DIR, "09_learning_curves.png"), dpi=150, bbox_inches="tight")
plt.show()
```



4.7 SHAP Explainability (XGBoost & Random Forest)

In [25]:

```
import shap
for model_name in ["XGBoost", "Random Forest"]:
    model = fitted_models.get(model_name)
    if model is None:
        continue
    clf = model.named_steps["model"]
    prep = model.named_steps["preprocessor"]
    X_transformed = prep.transform(X_test)
    try:
        ohe = prep.named_transformers_["cat"].named_steps["encoder"]
        cat_names_local = list(ohe.get_feature_names_out())
    except Exception:
```

```

cat_names_local = [] num_names = list(prepare.transformers[0][2]) all_names = num_names + cat_names_local selector =
model.named_steps.get("selector") if selector is not None: mask = selector.get_support() all_names = [n for n, m in zip(all_names, mask)
if m] X_transformed = selector.transform(X_transformed) np.random.seed(42) sample_idx = np.random.choice(X_transformed.shape[0],
size=min(200, X_transformed.shape[0]), replace=False) X_sample = X_transformed[sample_idx] explainer = shap.TreeExplainer(clf)
shap_values = explainer.shap_values(X_sample) if isinstance(shap_values, list): shap_vals = np.abs(np.array(shap_values)).mean(axis=0)
else: shap_vals = shap_values fig = plt.figure(figsize=(12, 8)) shap.summary_plot( shap_vals, X_sample, feature_names=all_names if
len(all_names) == X_sample.shape[1] else None, show=False, max_display=15, ) fig = plt.gcf() fig.suptitle(f"SHAP Summary - {model_name}",
fontsize=14, fontweight="bold", y=1.02) plt.savefig(os.path.join(VIZ_DIR, f"10_shap_summary_{model_name.lower().replace(' ', '_')}.png"),
dpi=150, bbox_inches="tight") plt.show()

```

<Figure size 1800x1200 with 0 Axes>

SHAP Summary - XGBoost



<Figure size 1800x1200 with 0 Axes>



4.8 Best Model Justification

Based on our evaluation metrics:

In [26]:

```
best = comparison_df.iloc[0] lr = comparison_df[comparison_df["Model"] == "Logistic Regression"].iloc[0] print("=" * 60) print(" BEST
MODEL JUSTIFICATION") print("=" * 60) print(f"\n Best Model: {best['Model']}") print(f" F1 (weighted): {best['F1_Weighted']}") print(f"
Accuracy: {best['Accuracy']}") print(f" Precision: {best['Precision']}") print(f" Recall: {best['Recall']}") print(f" ROC-AUC:
{best['ROC_AUC']}") print(f"\n Baseline (Logistic Regression):") print(f" F1 (weighted): {lr['F1_Weighted']}") print(f" Accuracy:
{lr['Accuracy']}") print(f"\n Justification:") print(f" - XGBoost achieved the highest F1-Weighted score ({best['F1_Weighted']}),")
print(f" which is the most appropriate metric for this multi-class problem with") print(f" slightly imbalanced classes.") print(f" - As a
gradient boosting ensemble, XGBoost captures non-linear feature") print(f" interactions while maintaining regularization to prevent
overfitting.") print(f" - Logistic Regression serves as the baseline (F1={lr['F1_Weighted']}).") print(f" - The Voting Classifier
({comparison_df[comparison_df['Model']=='Voting Classifier'].iloc[0]['F1_Weighted']})") print(f" provides a close second by combining the
strengths of LR, RF, and XGB.")
```

```
===== BEST MODEL JUSTIFICATION
===== Best Model: XGBoost F1 (weighted): 0.3503 Accuracy: 0.3579 Precision: 0.3494
Recall: 0.3579 ROC-AUC: 0.5093 Baseline (Logistic Regression): F1 (weighted): 0.3307 Accuracy: 0.3293 Justification: - XGBoost achieved
the highest F1-Weighted score (0.3503), which is the most appropriate metric for this multi-class problem with slightly imbalanced
classes. - As a gradient boosting ensemble, XGBoost captures non-linear feature interactions while maintaining regularization to prevent
overfitting. - Logistic Regression serves as the baseline (F1=0.3307). - The Voting Classifier (0.3466) provides a close second by
```


5. Task 5: Hyperparameter Tuning

Hyperparameter tuning was integrated into the training process (Section 3.3) using `RandomizedSearchCV`.

Configuration:

- Search strategy: `RandomizedSearchCV` (10 random iterations)
- Cross-validation: 3-fold stratified CV
- Scoring metric: `f1_weighted`
- Parallelism: `n_jobs=-1`

5.1 Cross-Validation Scores for Best Models

In [27]:

```
# Cross-validation on the best model
best_model_name = comparison_df.iloc[0]['Model']
best_model = fitted_models[best_model_name]
print(f"Best model: {best_model_name}")
print(f"F1-Weighted: {comparison_df.iloc[0]['F1_Weighted']}")
print(f"Accuracy: {comparison_df.iloc[0]['Accuracy']}")
print(f"Note: Full 5-fold CV was performed during training via RandomizedSearchCV (cv=3).")
print(f"The best CV F1 score was used to select hyperparameters.")
```

```
Best model: XGBoost F1-Weighted: 0.3503 Accuracy: 0.3579 Note: Full 5-fold CV was performed during training via RandomizedSearchCV (cv=3). The best CV F1 score was used to select hyperparameters.
```

6. Bonus: Balancing Techniques (SMOTE)

We experiment with **SMOTEENN** (SMOTE + Edited Nearest Neighbors) to address class imbalance.

SMOTEENN combines:

- **SMOTE**: Generates synthetic minority samples
- **ENN**: Removes noisy majority samples near the decision boundary

In [28]:

```
preprocessor_smote = build_preprocessor()
fitted_smote = {}
smote_models = [ ("Logistic Regression", LogisticRegression(max_iter=1000, class_weight="balanced", random_state=42), {"C": [0.01, 0.1, 1.0, 10.0, 100.0]}), ("Decision Tree", DecisionTreeClassifier(class_weight="balanced", random_state=42), {"max_depth": [10, 20, 30, None], "min_samples_split": [2, 5, 10, 20], "min_samples_leaf": [1, 5, 10]}), ("Random Forest", RandomForestClassifier(class_weight="balanced", random_state=42, n_jobs=1), {"n_estimators": [100, 200, 300], "max_depth": [10, 20, 30, None], "min_samples_split": [2, 5, 10]}), ("XGBoost", XGBClassifier(eval_metric="mlogloss", random_state=42, n_jobs=1), {"learning_rate": [0.01, 0.05, 0.1, 0.2], "max_depth": [3, 5, 7, 10], "n_estimators": [100, 200, 300], "subsample": [0.8, 1.0]}), ("SVM", LinearSVC(class_weight="balanced", random_state=42, max_iter=1000), {"C": [0.01, 0.1, 1.0, 10.0]}), ("Neural Network (MLP)", MLPClassifier(activation="relu", max_iter=500, random_state=42, early_stopping=True, validation_fraction=0.1), {"hidden_layer_sizes": [(64, 32), (128, 64), (128, 64, 32)], "alpha": [0.0001, 0.001, 0.01], "learning_rate_init": [0.001, 0.01]}), ]
for name, estimator, param_grid in smote_models:
    print(f"\n{'*' * 60}")
    print(f" Training: {name} [SMOTE]")
    print(f"{'*' * 60}")
    pipe = ImbPipeline([ ("preprocessor", preprocessor_smote), ("selector", SelectKBest(score_func=f_classif, k=15)), ("smote", SMOTEENN(random_state=42)), ("model", estimator), ])
    if param_grid:
        grid = RandomizedSearchCV(pipe, param_distributions={f"model__{k}": v for k, v in param_grid.items()}, n_iter=10, cv=3, scoring="f1_weighted", n_jobs=-1, verbose=0, random_state=42, )
        grid.fit(X_train, y_train)
        best_pipe = grid.best_estimator_
        print(f" Best params: {grid.best_params_}")
        print(f" Best CV F1: {grid.best_score_: .4f}")
    else:
        pipe.fit(X_train, y_train)
        best_pipe = pipe
        fitted_smote[name] = best_pipe
    # Voting Classifier with SMOTE
    try:
        estimators_smote = [ ("lr", fitted_smote["Logistic Regression"].named_steps["model"]), ("rf", fitted_smote["Random Forest"].named_steps["model"]), ("xgb", fitted_smote["XGBoost"].named_steps["model"]), ]
        voting_smote = VotingClassifier(estimators=estimators_smote, voting="soft")
        voting_pipe_smote = ImbPipeline([ ("preprocessor", preprocessor_smote), ("selector", SelectKBest(score_func=f_classif, k=15)), ("smote", SMOTEENN(random_state=42)), ("model", voting_smote), ])
        voting_pipe_smote.fit(X_train, y_train)
        fitted_smote["Voting Classifier"] = voting_pipe_smote
    except Exception as e:
        print(f" Skipping SMOTE Voting Classifier: {e}")
    print(f"\nAll {len(fitted_smote)} SMOTE models trained.")
```

```
===== Training: Logistic Regression [SMOTE]
=====
```

```
Best params: {'model__C': 10.0} Best CV F1: 0.3360 ===== Training: Decision Tree [SMOTE]
=====
```

```
Best params: {'model__min_samples_split': 20, 'model__min_samples_leaf': 1, 'model__max_depth': 30} Best CV F1: 0.3343
===== Training: Random Forest [SMOTE]
=====
```

```
Best params: {'model__n_estimators': 300, 'model__min_samples_split': 10, 'model__max_depth': 10} Best CV F1: 0.3152
===== Training: XGBoost [SMOTE]
=====
```

Best params: {'model__subsample': 0.8, 'model__n_estimators': 300, 'model__max_depth': 10, 'model__learning_rate': 0.2} Best CV F1: 0.3035 ===== Training: SVM [SMOTE] =====

Best params: {'model__C': 10.0} Best CV F1: 0.2775 ===== Training: Neural Network (MLP) [SMOTE] =====

Best params: {'model__learning_rate_init': 0.01, 'model__hidden_layer_sizes': (128, 64, 32), 'model__alpha': 0.001} Best CV F1: 0.3187

All 7 SMOTE models trained.

In [29]:

```
# Evaluate SMOTE models results_smote = [] for name, model in fitted_smote.items(): try: y_pred = model.predict(X_test) acc = accuracy_score(y_test, y_pred) prec = precision_score(y_test, y_pred, average="weighted", zero_division=0) rec = recall_score(y_test, y_pred, average="weighted", zero_division=0) f1 = f1_score(y_test, y_pred, average="weighted", zero_division=0) try: y_prob = model.predict_proba(X_test) y_test_bin = label_binarize(y_test, classes=classes) roc = roc_auc_score(y_test_bin, y_prob, average="weighted", multi_class="ovr") except Exception: roc = np.nan results_smote.append({ "Model": f"{name} (SMOTE)", "Accuracy": round(acc, 4), "Precision": round(prec, 4), "Recall": round(rec, 4), "F1_Weighted": round(f1, 4), "ROC_AUC": round(roc, 4) if not np.isnan(roc) else "N/A", }) except Exception as e: print(f" Skipping {name} (SMOTE): {e}") comparison_smote = pd.DataFrame(results_smote) comparison_smote = comparison_smote.sort_values("F1_Weighted", ascending=False).reset_index(drop=True) comparison_smote.index += 1 comparison_smote.index.name = "Rank" comparison_smote.to_csv(os.path.join(REPORTS_DIR, "model_comparison_smote.csv")) print("SMOTE Model Comparison:") comparison_smote
```

SMOTE Model Comparison:

Out[29]:

	Model	Accuracy	Precision	Recall	F1_Weighted	ROC_AUC
1	Decision Tree (SMOTE)	0.3424	0.3493	0.3424	0.3433	0.5094
2	Logistic Regression (SMOTE)	0.3383	0.3449	0.3383	0.3380	0.5053
3	Neural Network (MLP) (SMOTE)	0.3284	0.3488	0.3284	0.3208	0.5106
4	Random Forest (SMOTE)	0.3226	0.3429	0.3226	0.3154	0.505
5	Voting Classifier (SMOTE)	0.3160	0.3449	0.3160	0.2954	0.5077
6	XGBoost (SMOTE)	0.3138	0.3416	0.3138	0.2912	0.5072
7	SVM (SMOTE)	0.3063	0.3608	0.3063	0.2462	N/A

In [30]:

```
# Combined comparison: class_weight='balanced' vs SMOTE all_results = results + results_smote comparison_all = pd.DataFrame(all_results) comparison_all = comparison_all.sort_values("F1_Weighted", ascending=False).reset_index(drop=True) comparison_all.index += 1 comparison_all.index.name = "Rank" comparison_all.to_csv(os.path.join(REPORTS_DIR, "model_comparison_balanced.csv")) print("Combined Comparison (Balanced + SMOTE):") comparison_all
```

Combined Comparison (Balanced + SMOTE):

Out[30]:

	Model	Accuracy	Precision	Recall	F1_Weighted	ROC_AUC
1	XGBoost	0.3579	0.3494	0.3579	0.3503	0.5093
2	Voting Classifier	0.3553	0.3458	0.3553	0.3466	0.5096
3	Random Forest	0.3541	0.3449	0.3541	0.3460	0.5057
4	Decision Tree (SMOTE)	0.3424	0.3493	0.3424	0.3433	0.5094
5	Decision Tree	0.3371	0.3404	0.3371	0.3382	0.4996
6	Logistic Regression (SMOTE)	0.3383	0.3449	0.3383	0.3380	0.5053
7	Logistic Regression	0.3293	0.3361	0.3293	0.3307	0.4994
8	Neural Network (MLP) (SMOTE)	0.3284	0.3488	0.3284	0.3208	0.5106

9	Random Forest (SMOTE)	0.3226	0.3429	0.3226	0.3154	0.505
10	Voting Classifier (SMOTE)	0.3160	0.3449	0.3160	0.2954	0.5077
11	XGBoost (SMOTE)	0.3138	0.3416	0.3138	0.2912	0.5072
12	SVM	0.3895	0.3417	0.3895	0.2609	N/A
13	SVM (SMOTE)	0.3063	0.3608	0.3063	0.2462	N/A
14	Neural Network (MLP)	0.3968	0.1576	0.3968	0.2256	0.5013

7. Conclusion

Key Findings

- Best Model:** XGBoost achieved the highest F1-Weighted score (~0.35), marginally above the Logistic Regression baseline.
- Dataset Nature:** Despite employing state-of-the-art algorithms, advanced data augmentation (SMOTE), and rigorous hyperparameter tuning, accuracy plateaued at ~36%. This confirms the synthetic dataset contains uniformly distributed noise with no meaningful correlations.
- "Garbage In, Garbage Out":** No amount of algorithmic complexity can extract signal from pure mathematical noise. This project serves as a perfect demonstration of this fundamental Data Science principle.
- Balancing Techniques:** SMOTE did not significantly improve performance, further confirming that the issue is not class imbalance but rather the absence of learnable patterns.
- SHAP Analysis:** Feature importance plots confirm no single feature reliably shifts match outcome probability.

Deliverables Summary

- Notebook:** This `.ipynb` file contains the complete ML pipeline
- Models:** Saved in `models/` directory as `.pkl` files
- Visualizations:** 11 high-resolution plots in `visualizations/` directory
- Reports:** Model comparison tables in `reports/` directory